




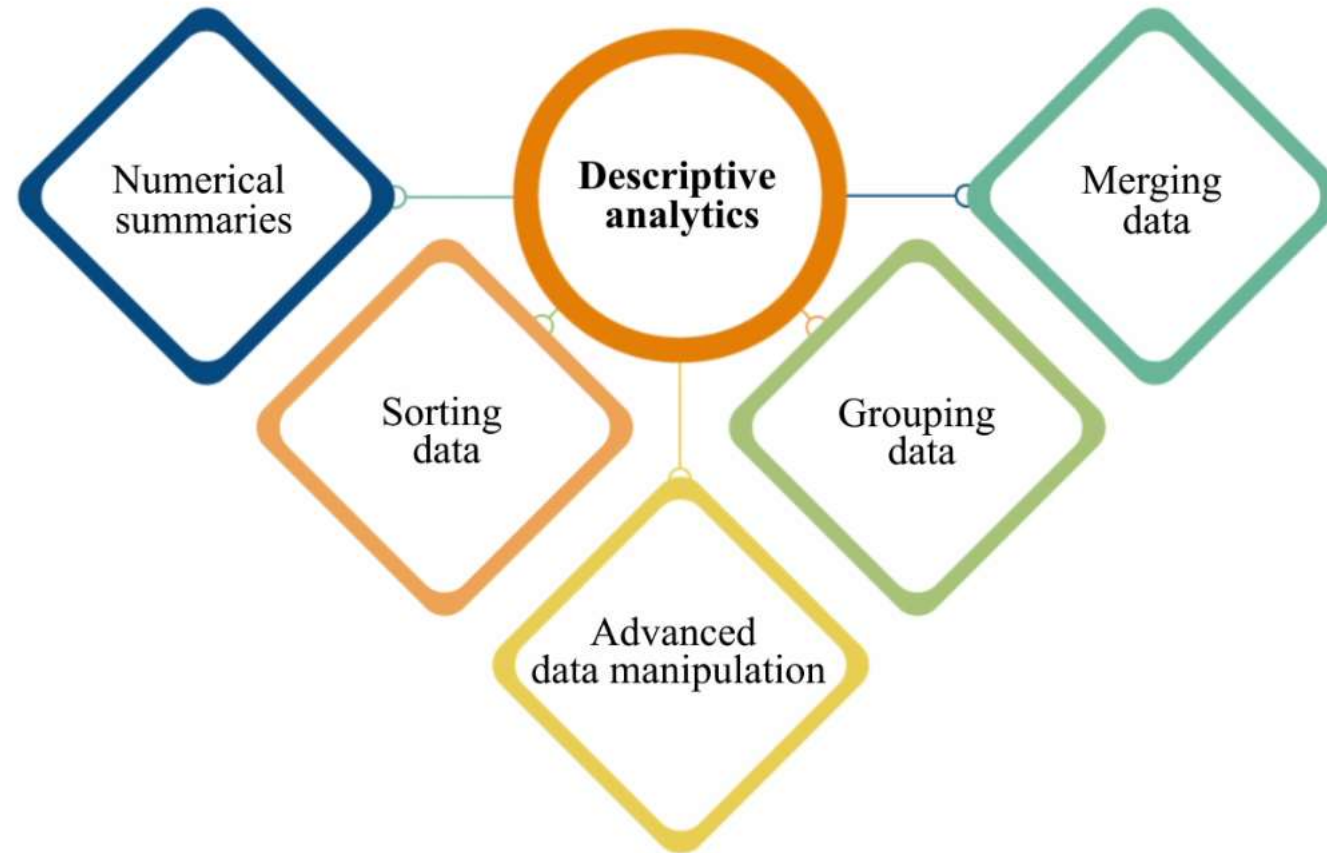
Python for Analytics: Week 7

Descriptive Analytics: Numerical Summary

The slide features a large blue triangle on the left side, pointing towards the bottom right. In the bottom left corner, there is a small orange triangle pointing towards the top right.

Descriptive Analytics with Numerical Summary: Overview

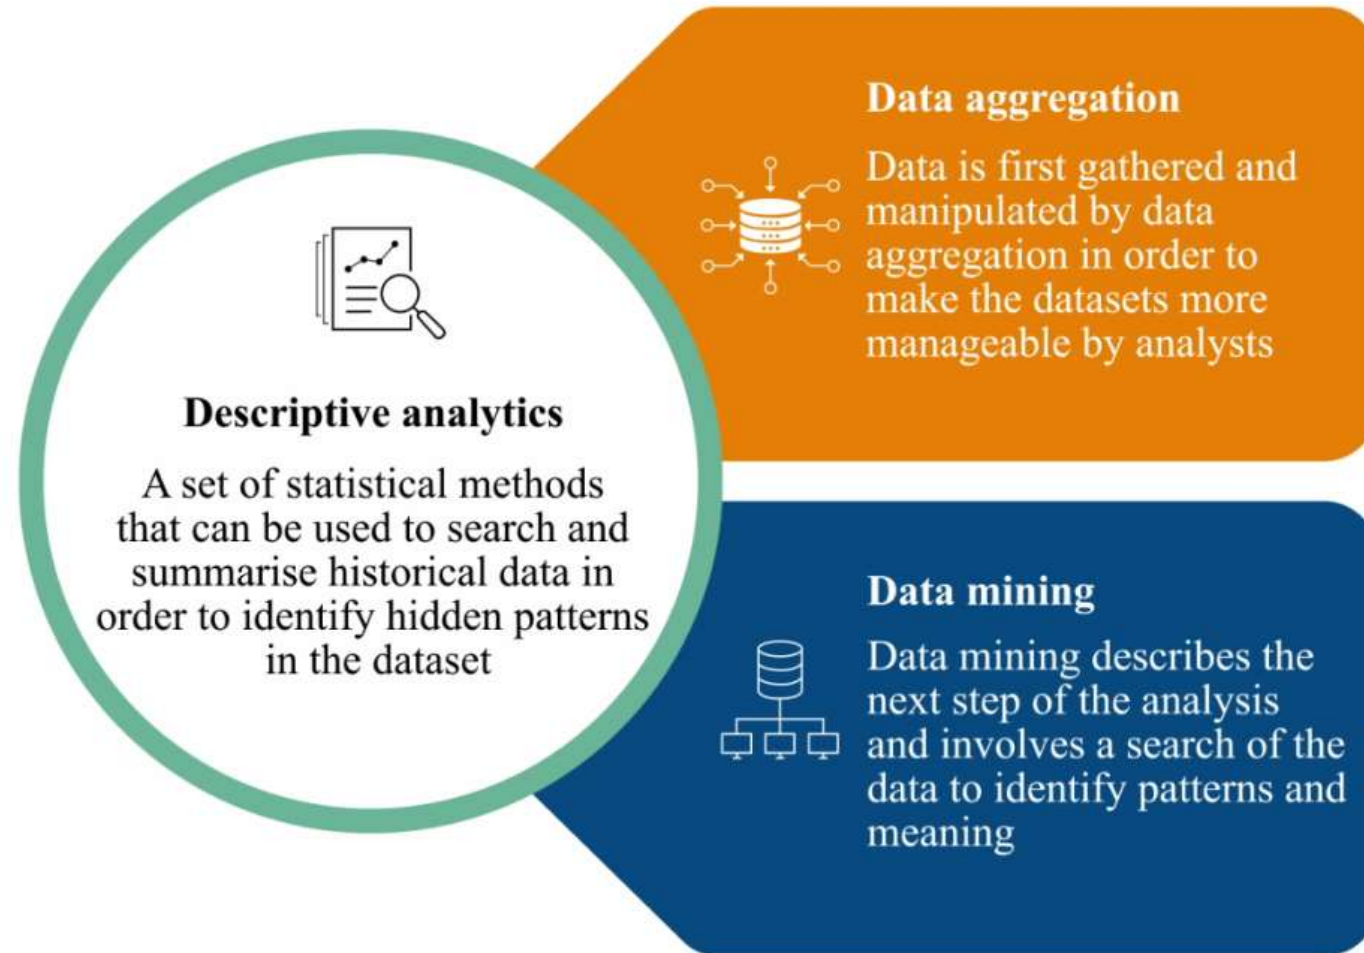
Module Overview





Introduction to Descriptive Analysis

Major Techniques in Descriptive Analysis



Reading Data in Python

```
In [38]: 1 data = pd.read_csv('wage.csv') # Read data from a file "wage.csv"
          2 data.head(6)                # Return the first 6 rows
```

Out[38]:

| | wage | educ | exper | female | married |
|---|------|------|-------|--------|---------|
| 0 | 3.10 | 11.0 | 2.0 | 1.0 | 0.0 |
| 1 | 3.24 | 12.0 | 22.0 | 1.0 | 1.0 |
| 2 | 3.00 | 11.0 | 2.0 | 0.0 | 0.0 |
| 3 | 6.00 | 8.0 | 44.0 | 0.0 | 1.0 |
| 4 | 5.30 | 12.0 | 7.0 | 0.0 | 1.0 |
| 5 | 8.75 | 16.0 | 9.0 | 0.0 | 1.0 |

| | wage | educ | exper | female | married |
|-----|-------|------|-------|--------|---------|
| 0 | 3.10 | 11.0 | 2.0 | 1.0 | 0.0 |
| 1 | 3.24 | 12.0 | 22.0 | 1.0 | 1.0 |
| 2 | 3.00 | 11.0 | 2.0 | 0.0 | 0.0 |
| 3 | 6.00 | 8.0 | 44.0 | 0.0 | 1.0 |
| 4 | 5.30 | 12.0 | 7.0 | 0.0 | 1.0 |
| ... | ... | ... | ... | ... | ... |
| 521 | 15.00 | 16.0 | 14.0 | 1.0 | 1.0 |
| 522 | 2.27 | 10.0 | 2.0 | 1.0 | 0.0 |
| 523 | 4.67 | 15.0 | 13.0 | 0.0 | 1.0 |
| 524 | 11.56 | 16.0 | 5.0 | 0.0 | 1.0 |
| 525 | 3.50 | 14.0 | 5.0 | 1.0 | 0.0 |

526 rows x 5 columns


Measures of Centre



Aspects of Centre Measures

Key centre majors are:


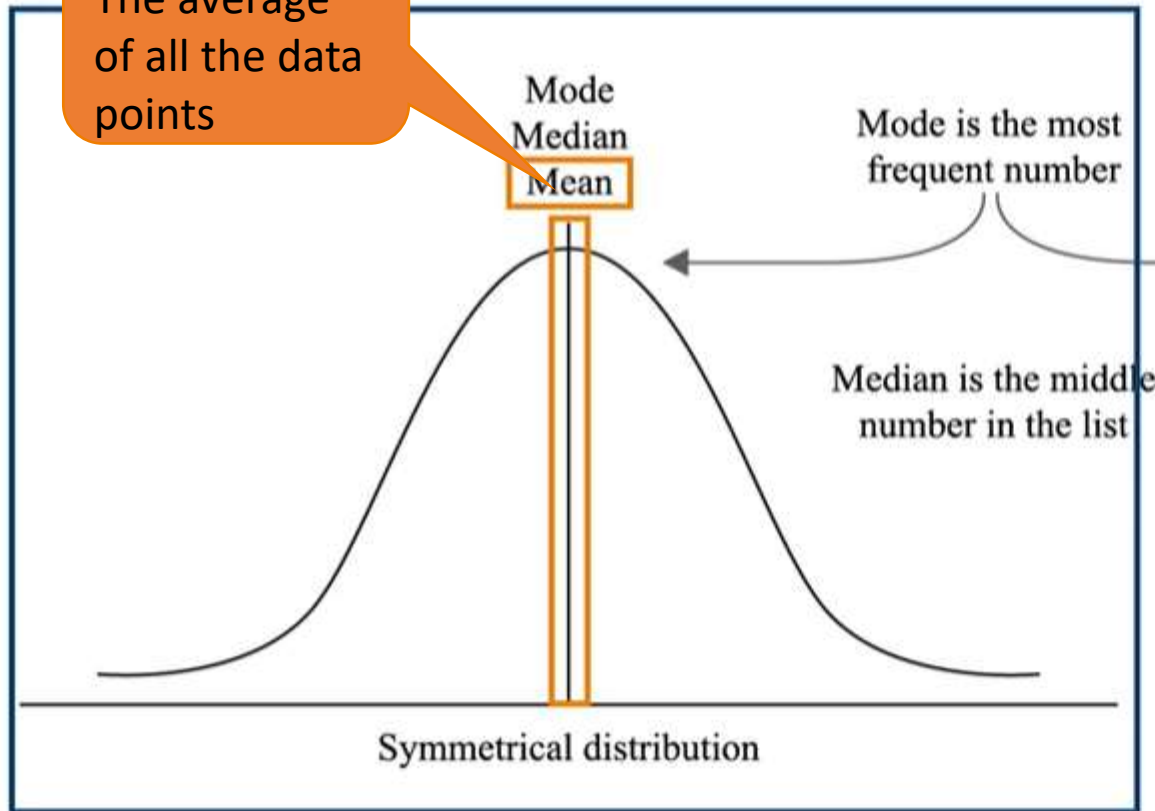
- Mean
- Median
- Mode



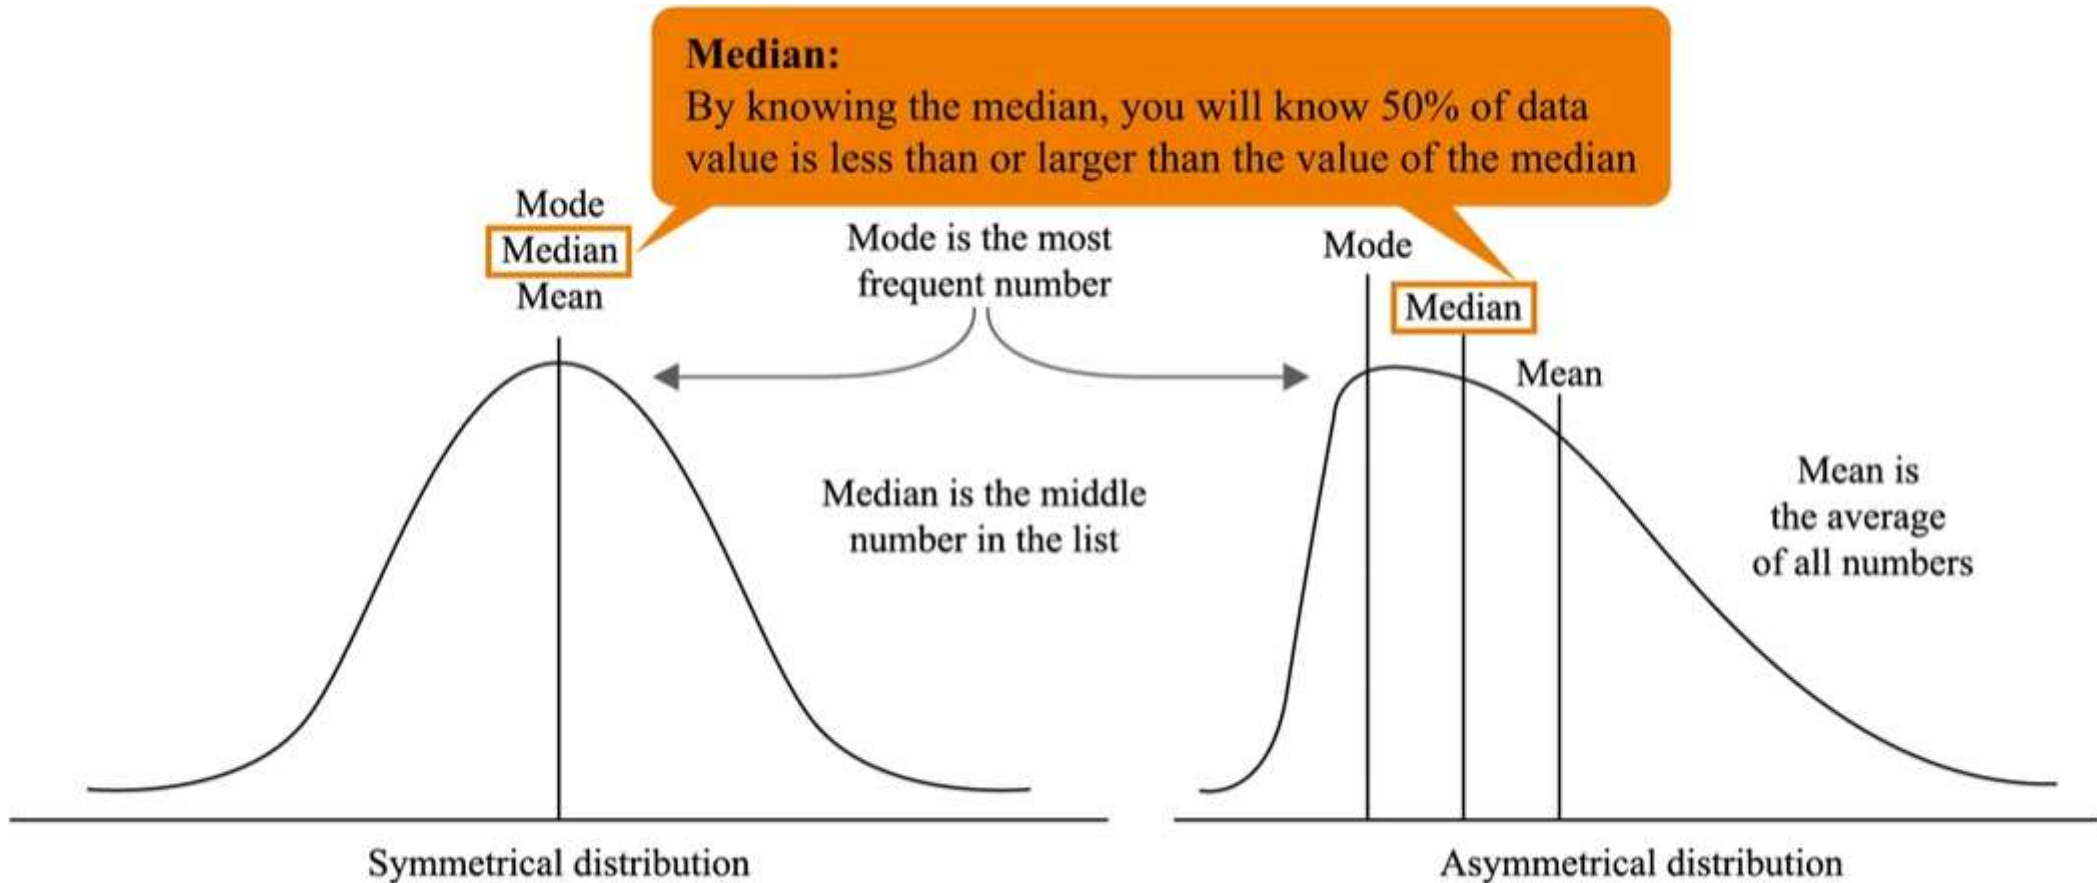
Aspects of Centre Measures

Mean:

The average of all the data points



Aspects of Centre Measures



Right-skewed distribution:

Most of the time, the mean value is larger than the median

Left-skewed distribution:

Most of the time, the median value is greater than the mean

Using the Mean Method: Pandas Series

Mean method calculates the summary statistics

In [39]:

```
1 print(data.mean())  
2 print(type(data.mean()))
```

```
Wage      5.896103  
Educ      12.562738  
Exper     17.017110  
Female     0.479087  
Married    0.608365  
dtype: float64
```

```
<class 'pandas.core.series.Series'>
```

| | Wage | Educ | Exper | Female | Married |
|-----|-------|------|-------|--------|---------|
| 0 | 3.10 | 11.0 | 2.0 | 1.0 | 0.0 |
| 1 | 3.24 | 12.0 | 22.0 | 1.0 | 1.0 |
| 2 | 3.00 | 11.0 | 2.0 | 0.0 | 0.0 |
| 3 | 6.00 | 8.0 | 44.0 | 0.0 | 1.0 |
| ... | ... | ... | ... | ... | ... |
| 522 | 2.27 | 10.0 | 2.0 | 1.0 | 0.0 |
| 523 | 4.67 | 15.0 | 13.0 | 0.0 | 1.0 |
| 524 | 11.56 | 16.0 | 5.0 | 0.0 | 1.0 |
| 525 | 3.50 | 14.0 | 5.0 | 1.0 | 0.0 |

data.mean()

| | | | | |
|-------|--------|--------|-------|-------|
| 5.896 | 12.563 | 17.017 | 0.479 | 0.608 |
|-------|--------|--------|-------|-------|

Using the Mean Method: Pandas Series

```
In [11]: print(wage_data.mean( ))  
         print(type(wage_data.mean( )))
```

```
Wage      5.896103  
Educ      12.562738  
Exper     17.017110  
Female    0.479087  
Married   0.608365  
dtype: float64
```

```
<class 'pandas.core.series.Series'>
```

$$\hat{p} = \frac{1}{N} \sum_{i=1}^N X_i$$

Total number
of cases

Number of females

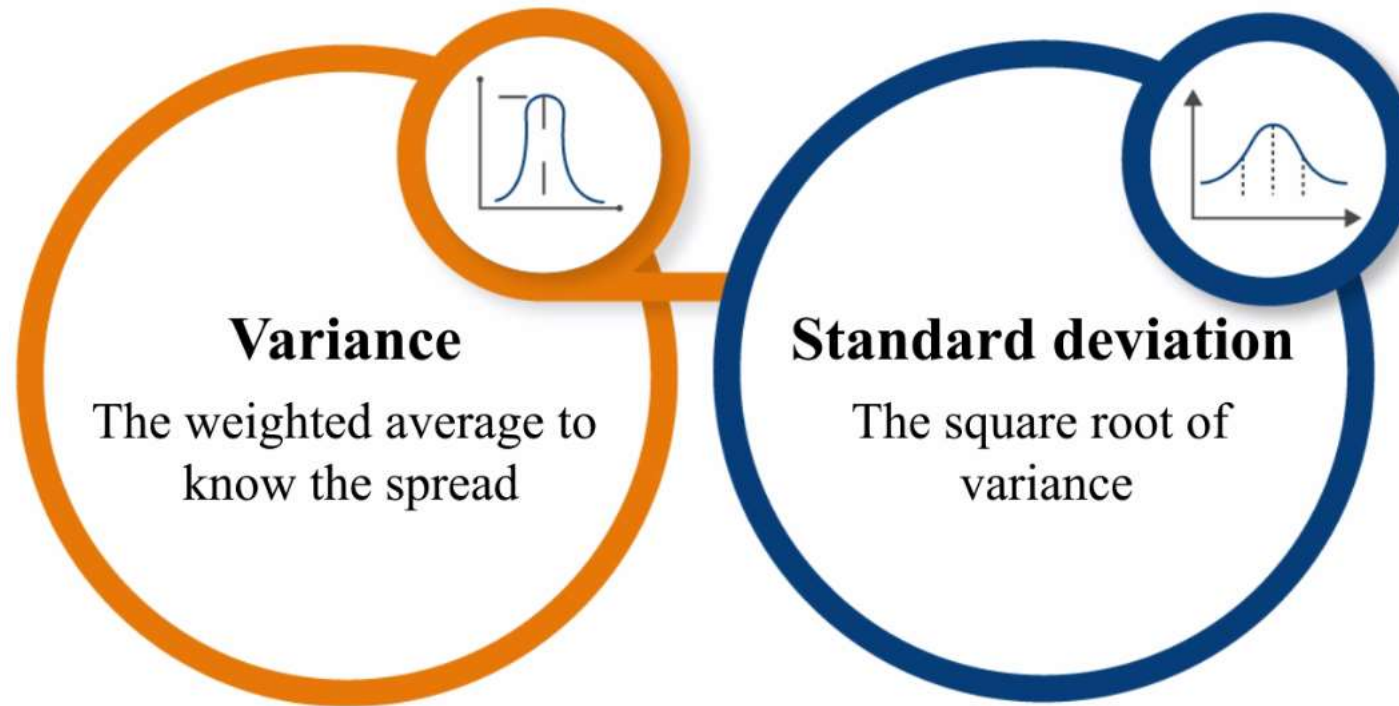
The dummy variable is:

$$X_i = \begin{cases} 1, & \text{if female} \\ 0, & \text{if male} \end{cases}$$


Measures of Variance and Extreme Points



Measures of Variation



Standard deviation




```
In [41]: 1 data.std()           # Sample Standard deviation of each column
```

```
Out[41]: wage           3.693086  
educ           2.769022  
exper          13.572160  
female          0.500038  
married          0.488580  
dtype: float64
```

Variance

Variance:

$$\hat{\sigma}^2 = \frac{1}{N-1} \sum_{i=1}^N (X_i - \bar{X})^2$$



```
In [42]: 1 data.var()           # Sample Variance of each column
```

```
Out[42]: wage      13.638884
educ       7.667485
exper     184.203516
female     0.250038
married    0.238711
dtype: float64
```

Extreme points

Minimum



Maximum



```
In [43]: 1 data.max()      # Maximum value of each column
```

```
Out[43]: wage      24.98  
educ      18.00  
exper     51.00  
female     1.00  
married     1.00  
dtype: float64
```

Describe Method



Describe Method

```
In [45]: 1 wage_summary = data.describe() # Obtain the key descriptive measures
        2 wage_summary                 # Display these measures as a table
```

Out[45]:


| | wage | educ | exper | female | married |
|--------------|------------|------------|------------|------------|------------|
| count | 526.000000 | 526.000000 | 526.000000 | 526.000000 | 526.000000 |
| mean | 5.896103 | 12.562738 | 17.01711 | 0.479087 | 0.608365 |
| std | 3.693086 | 2.769022 | 13.57216 | 0.500038 | 0.488580 |
| min | 0.530000 | 0.000000 | 1.00000 | 0.000000 | 0.000000 |
| 25% | 3.330000 | 12.000000 | 5.00000 | 0.000000 | 0.000000 |
| 50% | 4.650000 | 12.000000 | 13.50000 | 0.000000 | 1.000000 |
| 75% | 6.880000 | 14.000000 | 26.00000 | 1.000000 | 1.000000 |
| max | 24.980000 | 18.000000 | 51.00000 | 1.000000 | 1.000000 |

Q1
Q2
Q3

Quartiles: values
that divide a dataset
into quarters


Describe Method

Equal
distribution




(a) Uniform

Symmetric
distribution




(b) Bell shaped

Larger values on
the right tail



(c) Right skewed

Larger values on
the left tail



(d) Left skewed

The background features a large blue triangle on the left side, with a smaller orange triangle at the bottom-left corner.

Python Demo: Describe Method

Centres, Variations and Extreme Points

```
In [1]: import pandas as pd
```

```
In [2]: data_dict = {'wage': [3.10, 3.24, 3.00, 6.00, 5.30, 8.75],  
                    'educ': [11.0, 12.0, 11.0, 8.0, 12.0, 16.0],  
                    'exper': [2.0, 22.0, 2.0, 44.0, 7.0, 9.0],  
                    'female': [1.0, 1.0, 0.0, 0.0, 0.0, 0.0],  
                    'married': [0.0, 1.0, 0.0, 1.0, 1.0, 1.0]}  
  
data = pd.DataFrame(data_dict)    # DataFrame constructor  
data                               # Display the DataFrame
```

```
Out[2]:
```

| | wage | educ | exper | female | married |
|---|------|------|-------|--------|---------|
| 0 | 3.10 | 11.0 | 2.0 | 1.0 | 0.0 |
| 1 | 3.24 | 12.0 | 22.0 | 1.0 | 1.0 |
| 2 | 3.00 | 11.0 | 2.0 | 0.0 | 0.0 |
| 3 | 6.00 | 8.0 | 44.0 | 0.0 | 1.0 |
| 4 | 5.30 | 12.0 | 7.0 | 0.0 | 1.0 |
| 5 | 8.75 | 16.0 | 9.0 | 0.0 | 1.0 |

```
In [3]: data.mean()    # Mean value of each column
```

```
Out[3]: wage      4.898333  
educ      11.666667  
exper      14.333333  
female      0.333333  
married      0.666667  
dtype: float64
```

```
In [4]: type(data.mean())    # Show the data type of the results
```

```
Out[4]: pandas.core.series.Series
```

Centres, Variations and Extreme Points

```
In [4]: type(data.mean()) # Show the data type of the results|
```

```
Out[4]: pandas.core.series.Series
```

Use indexing to access
computed mean values

```
In [5]: data.median() # Median value of each column
```

```
Out[5]: wage      4.27  
educ      11.50  
exper      8.00  
female      0.00  
married      1.00  
dtype: float64
```

```
In [6]: type(data.median()) # Show the data type of the results|
```

```
Out[6]: pandas.core.series.Series
```

Please note that the mean value of each column is stored in a `pandas.Series`, where the index labels are variable names, rather than a sequence of integers.

Also notice that in the case of the 0-1 categorical variable "female", the mean value of 0.333 is the proportion of observations in the dataset labeled as "female". The same concept can be applied to the variable "married" as well.

Similarly, the measures of standard deviations and variances can be calculated by the corresponding methods.

Centres, Variations and Extreme Points

```
In [7]: data.std() | # Sample Standard deviation of each column
```

```
Out[7]: wage      2.271479  
educ      2.581989  
exper     16.280868  
female     0.516398  
married     0.516398  
dtype: float64
```

```
In [8]: data.var() # Sample Variance of each column
```

```
Out[8]: wage      5.159617  
educ      6.666667  
exper     265.066667  
female     0.266667  
married     0.266667  
dtype: float64
```

The maximum and minimum points in the dataset can also be found.

```
In [9]: data.max() # Maximum value of each column
```

```
Out[9]: wage      8.75  
educ     16.00  
exper     44.00  
female     1.00  
married     1.00  
dtype: float64
```

describe() Method

Method describe

For `pandas.DataFrame` and `pandas.Series`, the method `describe` is a convenient tool to summarize some key measures altogether.

```
In [11]: wage_summary = data.describe() # Obtain the key descriptive measures
wage_summary                                # Display these measures as a table
```

Out[11]:

| | wage | educ | exper | female | married |
|-------|----------|-----------|-----------|----------|----------|
| count | 6.000000 | 6.000000 | 6.000000 | 6.000000 | 6.000000 |
| mean | 4.898333 | 11.666667 | 14.333333 | 0.333333 | 0.666667 |
| std | 2.271479 | 2.581989 | 16.280868 | 0.516398 | 0.516398 |
| min | 3.000000 | 8.000000 | 2.000000 | 0.000000 | 0.000000 |
| 25% | 3.135000 | 11.000000 | 3.250000 | 0.000000 | 0.250000 |
| 50% | 4.270000 | 11.500000 | 8.000000 | 0.000000 | 1.000000 |
| 75% | 5.825000 | 12.000000 | 18.750000 | 0.750000 | 1.000000 |
| max | 8.750000 | 16.000000 | 44.000000 | 1.000000 | 1.000000 |

The variable `wage_summary` is a `pandas.DataFrame` table where the index labels are the names of the descriptive measures.

Note that rows `25%`, `50%`, and `75%` represent the first (Q1), second(Q2), and third quartiles(Q3), respectively. The value $Q3 - Q1$ is called the interquartile range (IQR).

Sorting Data



Sort numerical column in a DataFrame –

```
df.sort_values(by = 'col1', ascending = True):
```

Two input arguments

- **by:** Sort data by a specific column
- **ascending:** Sort data in ascending order or descending order by providing **true** value

Sort multiple columns at the same time –

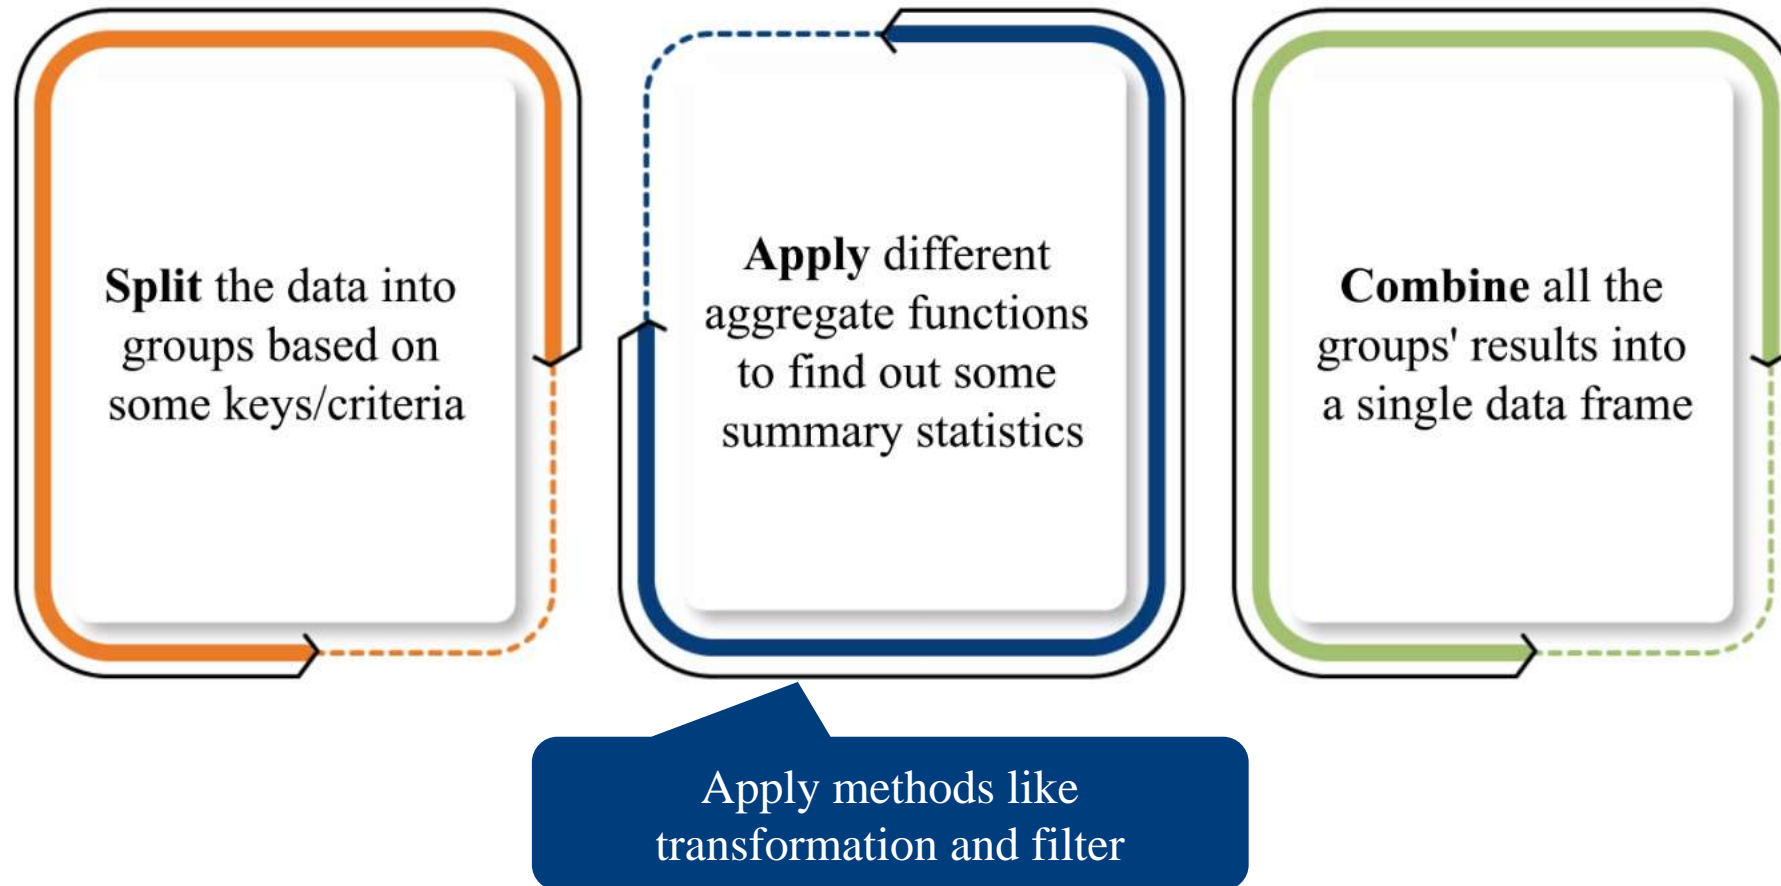
```
df.sort_values(by = ['col1', 'col2'],  
ascending = [True, False]):
```

- Provide a list of the columns in **by** argument
- Specify a Boolean list


The background features a large blue triangle on the left side, pointing towards the top-left corner. In the bottom-left corner, there is a small orange triangle pointing towards the bottom-left corner.

Advanced Data Manipulation Using Pandas: Introduction

Group-Wise Operators




Group-Wise Operators




Split by Groupby()


Cube A




Gender groupby



Year groupby




Class groupby



A
B
C
D


Split by Groupby()

Gender groupby




+

Year groupby




=

Gender-Year
groupby




Split by Groupby()

Gender-Year
groupby




Year-Class
groupby



Gender-Class
groupby



Gender – Year – Class
groupby



Groupby() Attributes

Downsize method

- Output the number of rows in each group
- ***size()***

Assess each groupby object

df_grouped['Score']:

Assess multiple columns for each group

df_grouped[['Class', 'Score']]:

Click to add text

Assess all groups

df_grouped.groups:

Assess specific group

df_grouped.get_group(('M', '2')):

Python Demo: Data Sorting



Sort Data

For numerical data, sorting in ascending order is important to help understand percentile values of the data. In Pandas, we have a useful method to sort values in a DataFrame or Series. The Pandas' method is `sort_values`

```
In [3]: data.sort_values(by = 'wage', ascending = False)
```

Out[3]:

| | wage | educ | exper | female | married |
|---|------|------|-------|--------|---------|
| 5 | 8.75 | 16.0 | 9.0 | 0.0 | 1.0 |
| 3 | 6.00 | 8.0 | 44.0 | 0.0 | 1.0 |
| 4 | 5.30 | 12.0 | 7.0 | 0.0 | 1.0 |
| 1 | 3.24 | 12.0 | 22.0 | 1.0 | 1.0 |
| 0 | 3.10 | 11.0 | 2.0 | 1.0 | 0.0 |
| 2 | 3.00 | 11.0 | 2.0 | 0.0 | 0.0 |

```
In [4]: data
```

Out[4]:

| | wage | educ | exper | female | married |
|---|------|------|-------|--------|---------|
| 0 | 3.10 | 11.0 | 2.0 | 1.0 | 0.0 |
| 1 | 3.24 | 12.0 | 22.0 | 1.0 | 1.0 |
| 2 | 3.00 | 11.0 | 2.0 | 0.0 | 0.0 |
| 3 | 6.00 | 8.0 | 44.0 | 0.0 | 1.0 |
| 4 | 5.30 | 12.0 | 7.0 | 0.0 | 1.0 |
| 5 | 8.75 | 16.0 | 9.0 | 0.0 | 1.0 |

Sort Data

```
In [6]: data.sort_values(by = 'wage', ascending = False, inplace = True)  
data
```

```
Out[6]:
```

| | wage | educ | exper | female | married |
|---|------|------|-------|--------|---------|
| 5 | 8.75 | 16.0 | 9.0 | 0.0 | 1.0 |
| 3 | 6.00 | 8.0 | 44.0 | 0.0 | 1.0 |
| 4 | 5.30 | 12.0 | 7.0 | 0.0 | 1.0 |
| 1 | 3.24 | 12.0 | 22.0 | 1.0 | 1.0 |
| 0 | 3.10 | 11.0 | 2.0 | 1.0 | 0.0 |
| 2 | 3.00 | 11.0 | 2.0 | 0.0 | 0.0 |

```
In [7]: data.sort_values(by = ['married', 'wage'], ascending = [False, True])
```

```
Out[7]:
```

| | wage | educ | exper | female | married |
|---|------|------|-------|--------|---------|
| 1 | 3.24 | 12.0 | 22.0 | 1.0 | 1.0 |
| 4 | 5.30 | 12.0 | 7.0 | 0.0 | 1.0 |
| 3 | 6.00 | 8.0 | 44.0 | 0.0 | 1.0 |
| 5 | 8.75 | 16.0 | 9.0 | 0.0 | 1.0 |
| 2 | 3.00 | 11.0 | 2.0 | 0.0 | 0.0 |
| 0 | 3.10 | 11.0 | 2.0 | 1.0 | 0.0 |

Sort Data

In [8]: data

Out[8]:

| | wage | educ | exper | female | married |
|---|------|------|-------|--------|---------|
| 5 | 8.75 | 16.0 | 9.0 | 0.0 | 1.0 |
| 3 | 6.00 | 8.0 | 44.0 | 0.0 | 1.0 |
| 4 | 5.30 | 12.0 | 7.0 | 0.0 | 1.0 |
| 1 | 3.24 | 12.0 | 22.0 | 1.0 | 1.0 |
| 0 | 3.10 | 11.0 | 2.0 | 1.0 | 0.0 |
| 2 | 3.00 | 11.0 | 2.0 | 0.0 | 0.0 |

In [9]: data.sort_values(by = ['married','wage'], ascending = [False,True], inplace = True)

In [10]: data

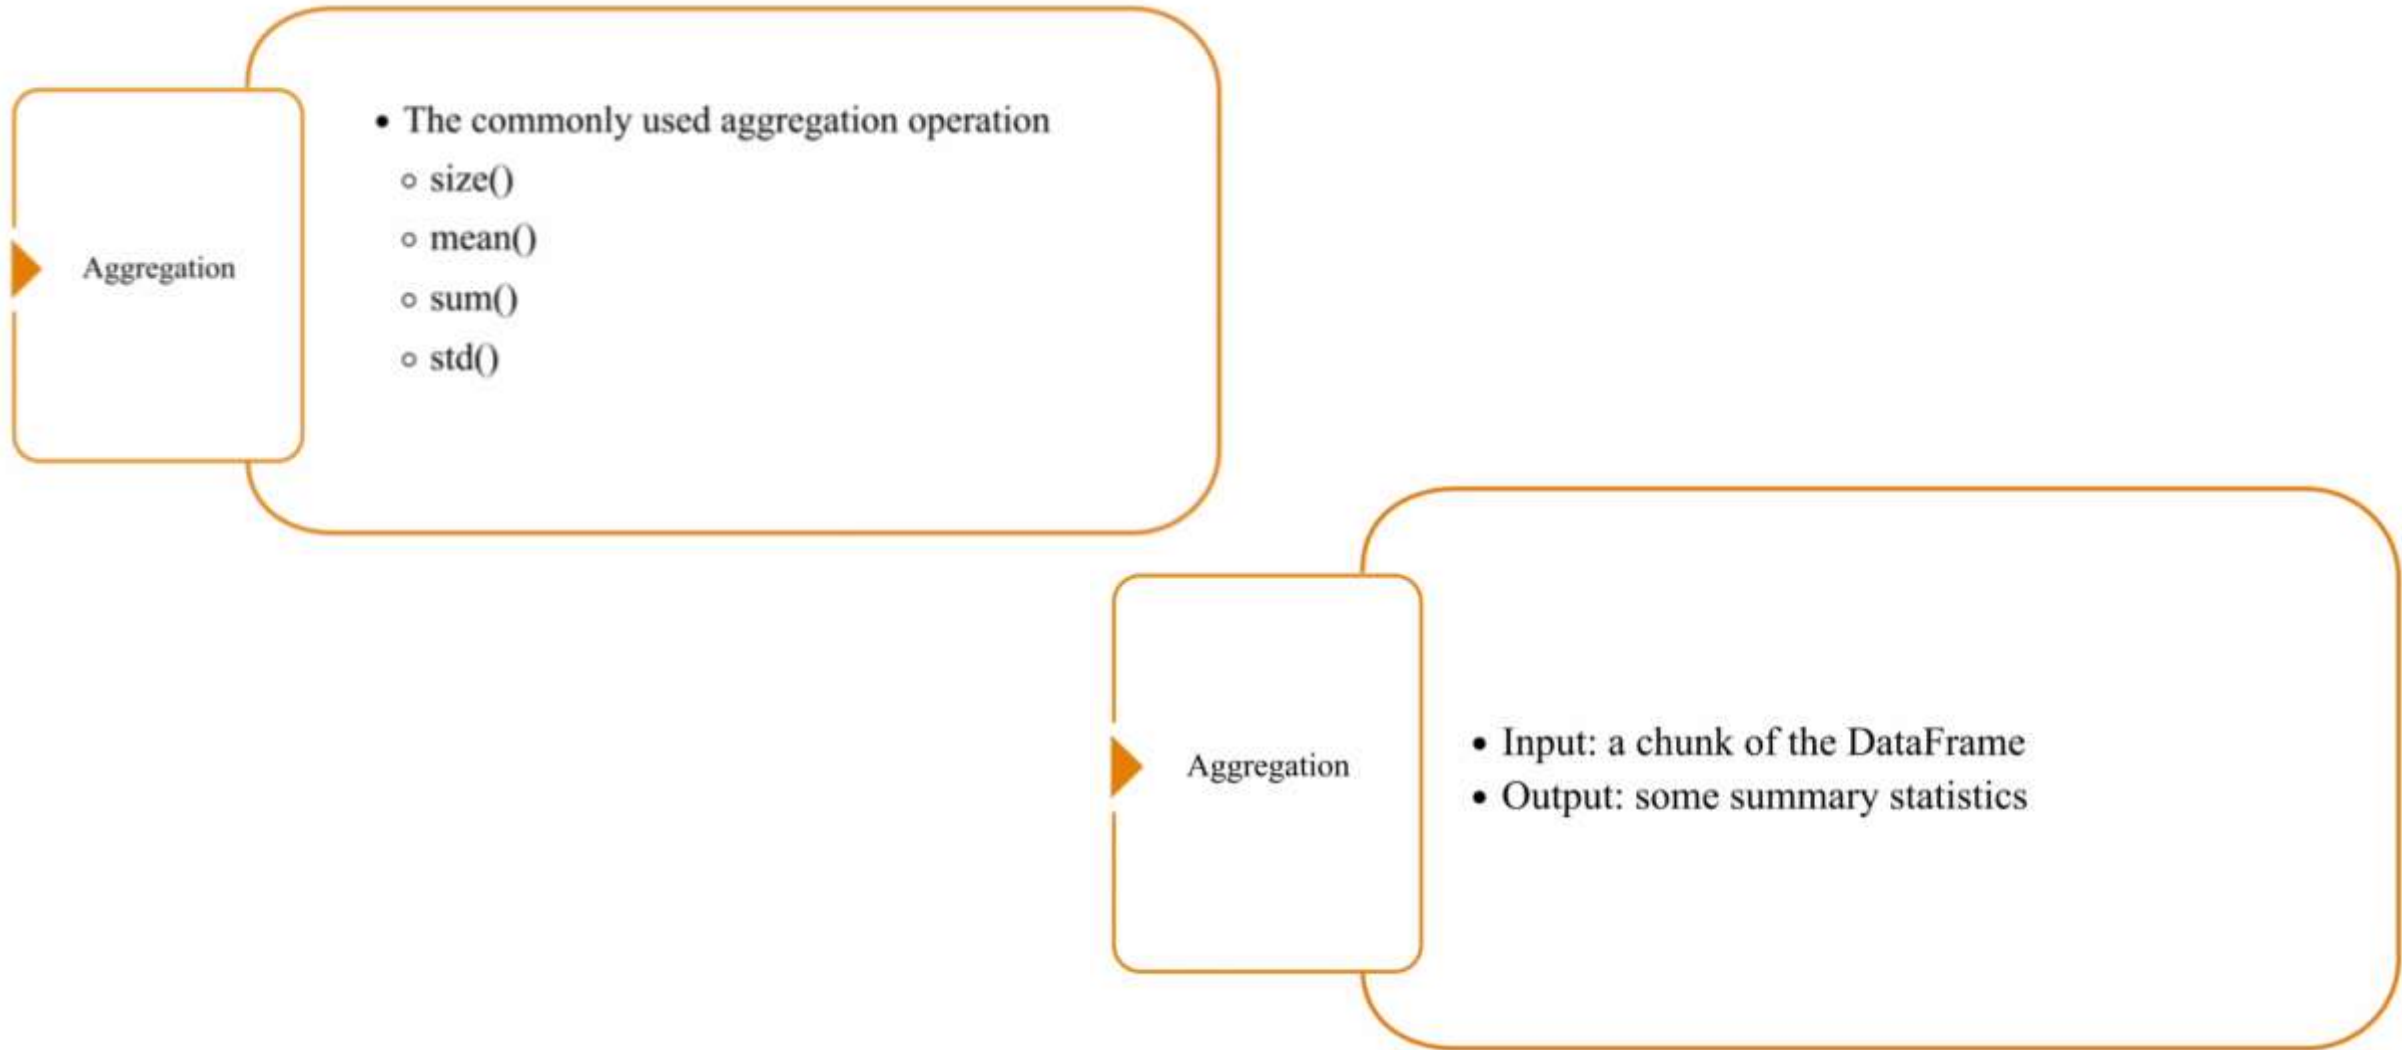
Out[10]:

| | wage | educ | exper | female | married |
|---|------|------|-------|--------|---------|
| 1 | 3.24 | 12.0 | 22.0 | 1.0 | 1.0 |
| 4 | 5.30 | 12.0 | 7.0 | 0.0 | 1.0 |
| 3 | 6.00 | 8.0 | 44.0 | 0.0 | 1.0 |
| 5 | 8.75 | 16.0 | 9.0 | 0.0 | 1.0 |
| 2 | 3.00 | 11.0 | 2.0 | 0.0 | 0.0 |
| 0 | 3.10 | 11.0 | 2.0 | 1.0 | 0.0 |

A large blue triangle occupies the left and bottom portions of the slide, pointing towards the bottom right corner.A small orange triangle is located in the bottom left corner of the slide.

Group-wise Operations: Aggregation, Transformation and Filtration

Group-wise Operations



Lambda Function

Special function in
Pandas

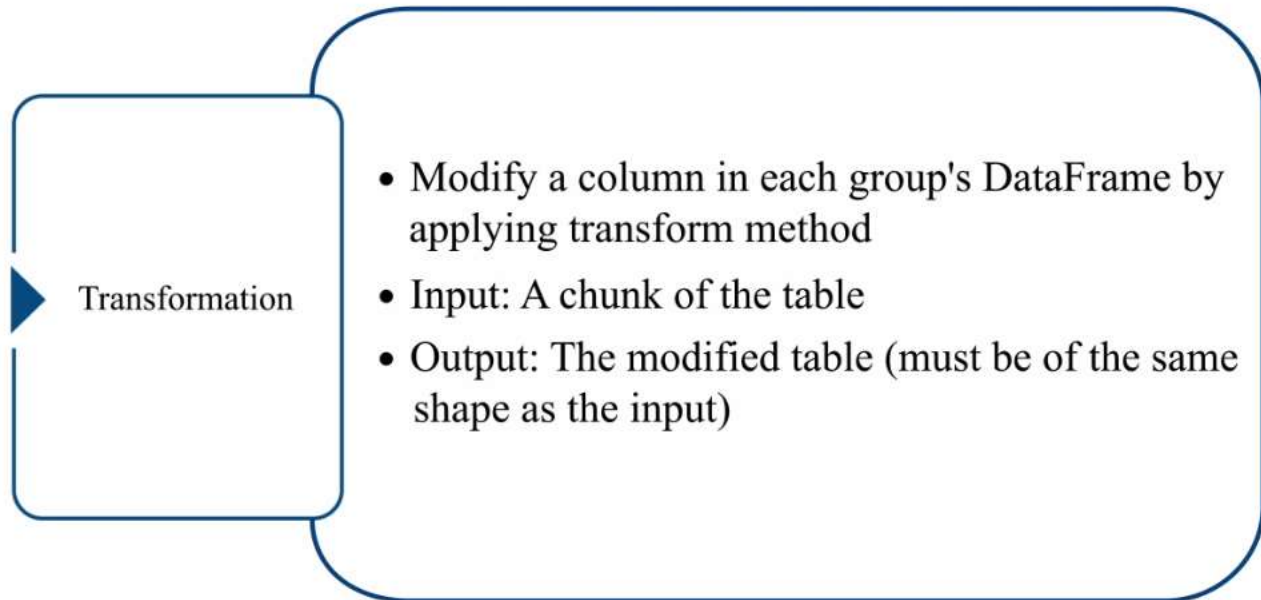
Can be viewed as a
single-use function
or anonymous
function

Contradicts tradition
function

- Does not have a function name
- Can have multiple input arguments
- Only has one-line expression as the return value

Example: `a = lambda x, y: x+y`

- Key word: '**lambda**' to start a Lambda Function
- Input arguments: **x** and **y** (Lambda Function can take any number of arguments)
- One-line function body and return value: **x+y**
- Type: **a(3,5)**



Example: Group-wise normalisation

- `zscore = lambda x: (x - x.mean()) / x.std()`
- `data.groupby(['Class'])[['Score']].transform(zscore)`

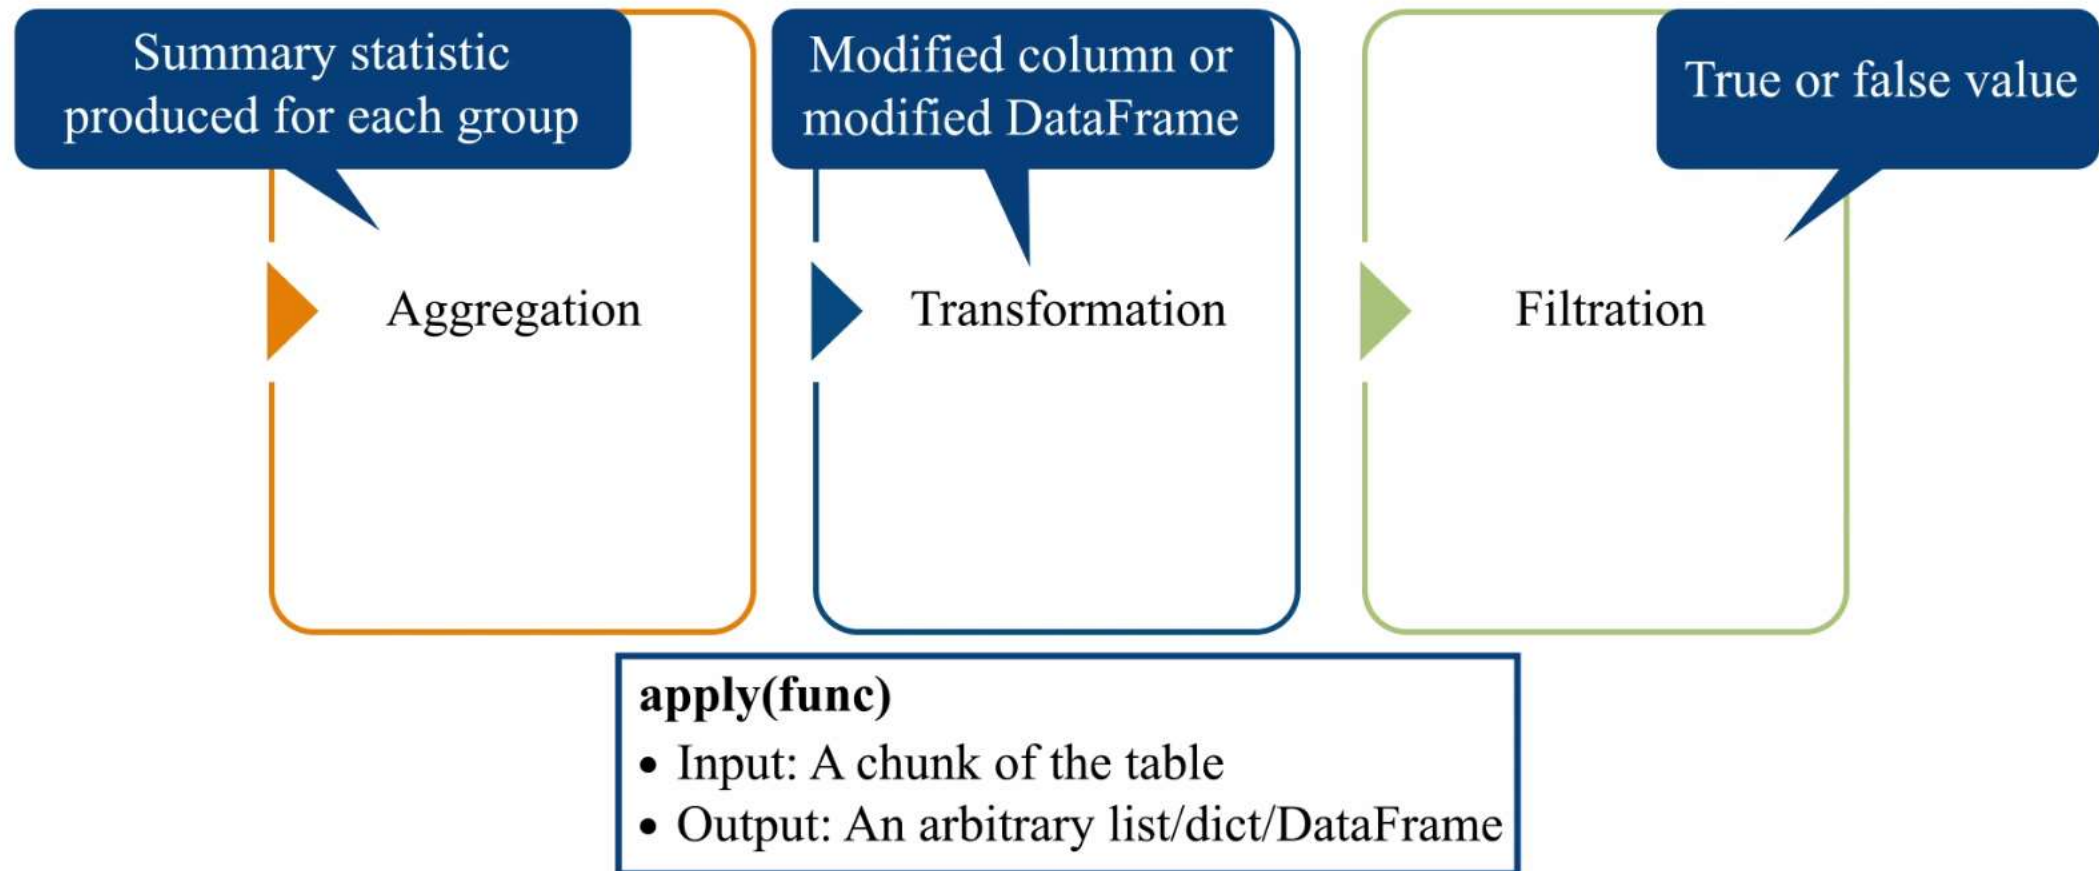
Filtration

- Filter groups by filter(func)
- Input: A chunk of the table
- Output: True/False

Example: Select those classes with a class average score > 80

- `high_class_avg = lambda x: x['Score'].mean() > 80`
- `data.groupby('Class').filter(high_class_avg)`

Group-wise Operators





Python Demo: Group-Wise Operations

Group-Wise Operations

Background: In the following, we will use Singapore's 4-Digits data set to exemplify the power of group-wise operations using Pandas. The 4-Digits (abbreviation: 4-D) is a lottery in Singapore. Individuals play by choosing any number from 0000 to 9999. Then, twenty-three winning numbers are drawn each time. If one of the numbers matches the one that the player has bought, a prize is won. A draw is conducted to select these winning numbers. 4-Digits is a fixed-odds game. There are five prize categories: **1st Prize**, **2nd Prize**, **3rd Prize**, **Starter Prizes** and **Consolation Prizes**. We would like to know prize-specific summary of the 4-D lottery.

Load data

The following data file "4D_results_long.csv" has been generated. First, please load data into Python.

```
In [11]: data = pd.read_csv('4D_results_long.csv')
```

```
In [ ]: data.head(10)
```

```
In [ ]: data.shape
```

Column Summary

value_counts()

```
In [ ]: # count the instances of each prize type  
data['prize_type'].value_counts()
```

```
In [ ]: # show the Top 10 most frequent winning number in the dataset  
data['number'].value_counts().head(10)
```

describe()

```
In [16]: # summary statistics of all the winning numbers  
data['number'].describe()
```

```
Out[16]: count    10810.000000  
         mean      4957.073636  
         std       2905.037198  
         min         0.000000  
         25%      2400.250000  
         50%      4964.000000  
         75%      7479.750000  
         max      9999.000000  
         Name: number, dtype: float64
```

```
In [ ]: # summary of a categorical variable  
data['weekday'].describe()
```

GroupBy

```
In [18]: data_by_prizetype = data.groupby(['prize_type'])
```

```
In [ ]: data_by_prizetype
```

```
In [20]: data_by_prizetype.size()
```

```
Out[20]: prize_type
1st          470
2nd          470
3rd          470
consolation  4700
starter      4700
dtype: int64
```

```
In [ ]:
```

```
In [ ]: #DataFrame column selection in GroupBy
data_by_prizetype['number']
```

```
In [ ]: data_by_prizetype[['number']]
```

```
In [ ]:
```

```
In [ ]: data_by_prizetype.groups
```

```
In [ ]: # Selecting a group
data_by_prizetype.get_group('2nd').head(10)
```

GroupBy

```
In [23]: data_by_prizetype.groups
```

```
In [ ]: # Selecting a group  
data_by_prizetype.get_group('2nd').head(10)
```

```
In [ ]:
```

```
In [ ]: # Iterating through groups  
for name, group in data_by_prizetype:  
    print(name)  
    print(group.shape)  
    print(type(group))
```

Aggregation

Mean of winning numbers for each prizetype

```
In [ ]: # as series  
data_by_prizetype['number'].mean()
```

```
In [ ]: # as dataframe with the group variable as index  
data_by_prizetype[['number']].mean()
```

```
In [ ]: # as dataframe without index  
data_by_prizetype[['number']].mean().reset_index()
```

```
In [ ]: # avoid index when grouping  
data.groupby('prize_type', as_index=False)[['number']].mean()
```

Count of prize_type for each draw_no

```
In [ ]: data.groupby('draw_no')['prize_type'].value_counts()
```

Aggregation by agg()

Applying multiple functions at once

```
In [ ]: import numpy as np
# Min/Median/Max of number for each weekday-prize_type combination
# group by 'weekday' and 'prize_type' and extract 'number' column
weekday_prize_gpby = data.groupby(['weekday', 'prize_type'])
```

```
In [ ]: weekday_prize_gpby
```

```
In [ ]: weekday_prize_gpby[['number']].agg([np.min, np.median, np.max])
```

Applying different functions to different columns

```
In [ ]: weekday_prize_gpby.agg({'number': np.mean, 'date': np.max})
```

agg with a customized function

```
In [36]: year_gpby = data.groupby('year')
```

```
In [ ]: year_gpby['number'].agg(lambda x: sum(x > 9900))
```

```
In [ ]: # count how many times that the prize number > 9900 for each year
year_gpby[['number']].agg(lambda x: sum(x > 9900))
```


Aggregation by agg()

Transformation

Normalize number by year

```
In [ ]: year_gpby[['number']].transform(lambda x: (x - x.mean())/x.std())
```

Filtration

Focus on first prizes, select data only in weeks with 3 draws

```
In [ ]: filt_1st = (data['prize'] == 'first_prize')
data_1st = data.loc[filt_1st,:].copy()      # Create a new data set with 1st prize data only
```

```
In [ ]: year_wkno_1st_gpby = data_1st.groupby(['year', 'week_no'])
```

```
In [ ]: year_wkno_1st_gpby.filter(lambda x: len(x) > 2)
```

Apply

For each year and each prize_type, find the draw with smallest number and the one with the largest number

```
In [ ]: def draw_min_max(x):
    i = x['number'].idxmin()
    j = x['number'].idxmax()
    return pd.concat([x.loc[i, ['draw_no', 'number']], x.loc[j, ['draw_no', 'number']]])
```

```
In [ ]: data.groupby(['year', 'prize_type']).apply(draw_min_max)
```

The slide features a large blue triangle on the left side, pointing towards the top-left corner. In the bottom-left corner, there is a small orange triangle pointing towards the bottom-left corner.

Joining and Merging Data: Introduction

Introduction to Joining and Merging Data

- Combine information from different sources
- Form a complete dataset
- Apply the dataset to answer business or analytics questions

Pandas methods to combine datasets

- Concat
- Merge

Concat and Merge



Concat Method

- Concat(): Combine DataFrames row-wise or column-wise by index
- Concat method is versatile
- One-to-one relationship only

Simple Concatenation: Example

df1

| | A | B | C | D |
|---|----|----|----|----|
| 0 | A0 | B0 | C0 | D0 |
| 1 | A1 | B1 | C1 | D1 |
| 2 | A2 | B2 | C2 | D2 |
| 3 | A3 | B3 | C3 | D3 |

df2

| | A | B | C | D |
|---|----|----|----|----|
| 4 | A4 | B4 | C4 | D4 |
| 5 | A5 | B5 | C5 | D5 |
| 6 | A6 | B6 | C6 | D6 |
| 7 | A7 | B7 | C7 | D7 |

df3

| | A | B | C | D |
|----|-----|-----|-----|-----|
| 8 | A8 | B8 | C8 | D8 |
| 9 | A9 | B9 | C9 | D9 |
| 10 | A10 | B10 | C10 | D10 |
| 11 | A11 | B11 | C11 | D11 |

`pd.concat([df1, df2, df3])`



Result

| | A | B | C | D |
|----|-----|-----|-----|-----|
| 0 | A0 | B0 | C0 | D0 |
| 1 | A1 | B1 | C1 | D1 |
| 2 | A2 | B2 | C2 | D2 |
| 3 | A3 | B3 | C3 | D3 |
| 4 | A4 | B4 | C4 | D4 |
| 5 | A5 | B5 | C5 | D5 |
| 6 | A6 | B6 | C6 | D6 |
| 7 | A7 | B7 | C7 | D7 |
| 8 | A8 | B8 | C8 | D8 |
| 9 | A9 | B9 | C9 | D9 |
| 10 | A10 | B10 | C10 | D10 |
| 11 | A11 | B11 | C11 | D11 |

Column-Wise Stacking: Example

| df1 | | | | | df4 | | | |
|-----|----|----|----|----|-----|----|----|----|
| | A | B | C | D | | B | D | F |
| 0 | A0 | B0 | C0 | D0 | 2 | B2 | D2 | F2 |
| 1 | A1 | B1 | C1 | D1 | 3 | B3 | D3 | F3 |
| 2 | A2 | B2 | C2 | D2 | 6 | B6 | D6 | F6 |
| 3 | A3 | B3 | C3 | D3 | 7 | B7 | D7 | F7 |

| Result | | | | | | | |
|--------|-----|-----|-----|-----|-----|-----|-----|
| | A | B | C | D | B | D | F |
| 0 | A0 | B0 | C0 | D0 | NaN | NaN | NaN |
| 1 | A1 | B1 | C1 | D1 | NaN | NaN | NaN |
| 2 | A2 | B2 | C2 | D2 | B2 | D2 | F2 |
| 3 | A3 | B3 | C3 | D3 | B3 | D3 | F3 |
| 6 | NaN | NaN | NaN | NaN | B6 | D6 | F6 |
| 7 | NaN | NaN | NaN | NaN | B7 | D7 | F7 |

Inner Join: Example

Inner join: Intersection of two set of indices

```
pd.concat([df1, df4], axis=1, join='inner')
```

| df1 | | | | |
|-----|----|----|----|----|
| | A | B | C | D |
| 0 | A0 | B0 | C0 | D0 |
| 1 | A1 | B1 | C1 | D1 |
| 2 | A2 | B2 | C2 | D2 |
| 3 | A3 | B3 | C3 | D3 |

| df4 | | | |
|-----|----|----|----|
| | B | D | F |
| 2 | B2 | D2 | F2 |
| 3 | B3 | D3 | F3 |
| 6 | B6 | D6 | F6 |
| 7 | B7 | D7 | F7 |

| Result | | | | | | | |
|--------|----|----|----|----|----|----|----|
| | A | B | C | D | B | D | F |
| 2 | A2 | B2 | C2 | D2 | B2 | D2 | F2 |
| 3 | A3 | B3 | C3 | D3 | B3 | D3 | F3 |

Merge Method

- Merge(): Column-wise stacking only (i.e., axis=1)
- One-to-many and many-to-many on top of one-to-one
- Align by multiple column (instead of by an index)
- Reference: <http://pandas.pydata.org/pandas-docs/stable/merging.html>

Merge Method: Example

| left | | | | right | | | | Result | | | | | |
|------|-----|----|----|-------|-----|----|----|--------|-----|----|----|----|----|
| | key | A | B | | key | C | D | | key | A | B | C | D |
| 0 | K0 | A0 | B0 | 0 | K0 | C0 | D0 | 0 | K0 | A0 | B0 | C0 | D0 |
| 1 | K1 | A1 | B1 | 1 | K1 | C1 | D1 | 1 | K1 | A1 | B1 | C1 | D1 |
| 2 | K2 | A2 | B2 | 2 | K2 | C2 | D2 | 2 | K2 | A2 | B2 | C2 | D2 |
| 3 | K3 | A3 | B3 | 3 | K3 | C3 | D3 | 3 | K3 | A3 | B3 | C3 | D3 |

Data Merge with Multiple Keys

| left | | | | | right | | | | |
|------|------|------|----|----|-------|------|------|----|----|
| | key1 | key2 | A | B | | key1 | key2 | C | D |
| 0 | K0 | K0 | A0 | B0 | 0 | K0 | K0 | C0 | D0 |
| 1 | K0 | K1 | A1 | B1 | 1 | K1 | K0 | C1 | D1 |
| 2 | K1 | K0 | A2 | B2 | 2 | K1 | K0 | C2 | D2 |
| 3 | K2 | K1 | A3 | B3 | 3 | K2 | K0 | C3 | D3 |

| Result | | | | | | |
|--------|------|------|----|----|----|----|
| | key1 | key2 | A | B | C | D |
| 0 | K0 | K0 | A0 | B0 | C0 | D0 |
| 1 | K1 | K0 | A2 | B2 | C1 | D1 |
| 2 | K1 | K0 | A2 | B2 | C2 | D2 |

Outer Join Method: Example

Outer join: **Union** of the keys, missing values filled with NaN

```
pd.merge(left, right, on=['key1', 'key2'], how='outer')
```

| left | | | | | right | | | | | Result | | | | | | |
|------|------|------|----|----|-------|------|------|----|----|--------|------|------|-----|-----|-----|-----|
| | key1 | key2 | A | B | | key1 | key2 | C | D | | key1 | key2 | A | B | C | D |
| 0 | K0 | K0 | A0 | B0 | 0 | K0 | K0 | C0 | D0 | 0 | K0 | K0 | A0 | B0 | C0 | D0 |
| 1 | K0 | K1 | A1 | B1 | 1 | K1 | K0 | C1 | D1 | 1 | K0 | K1 | A1 | B1 | NaN | NaN |
| 2 | K1 | K0 | A2 | B2 | 2 | K1 | K0 | C2 | D2 | 2 | K1 | K0 | A2 | B2 | C1 | D1 |
| 3 | K2 | K1 | A3 | B3 | 3 | K2 | K0 | C3 | D3 | 3 | K1 | K0 | A2 | B2 | C2 | D2 |
| | | | | | | | | | | 4 | K2 | K1 | A3 | B3 | NaN | NaN |
| | | | | | | | | | | 5 | K2 | K0 | NaN | NaN | C3 | D3 |

Merge Method: Outer Join Left Indices

| left | | | | | right | | | | |
|------|------|------|----|----|-------|------|------|----|----|
| | key1 | key2 | A | B | | key1 | key2 | C | D |
| 0 | K0 | K0 | A0 | B0 | 0 | K0 | K0 | C0 | D0 |
| 1 | K0 | K1 | A1 | B1 | 1 | K1 | K0 | C1 | D1 |
| 2 | K1 | K0 | A2 | B2 | 2 | K1 | K0 | C2 | D2 |
| 3 | K2 | K1 | A3 | B3 | 3 | K2 | K0 | C3 | D3 |

| Result | | | | | | |
|--------|------|------|-----|-----|----|----|
| | key1 | key2 | A | B | C | D |
| 0 | K0 | K0 | A0 | B0 | C0 | D0 |
| 1 | K1 | K0 | A2 | B2 | C1 | D1 |
| 2 | K1 | K0 | A2 | B2 | C2 | D2 |
| 3 | K2 | K0 | NaN | NaN | C3 | D3 |

Python Demo: Concat and Merge



Data Merge in Pandas

Background: In the following, we will use Google Public Data to illustrate Pandas' `concat` and `merge` methods. Please load in some data files first. We want to create a complete data set that can be used to explore the relationship between **fertility rate** and a country's **GDP per capita**. At the same time, two other demographic factors are also included: **household expenditure** and **population density of the country**. We will not do the detailed analysis here; instead, we just demonstrate how to prepare the required data using data concatenation and merging.

```
In [ ]: import pandas as pd
        data_fer = pd.read_csv("GooglePublicData_fer.csv")
        data_gdp = pd.read_csv("GooglePublicData_gdp.csv")
        data_exp = pd.read_csv("GooglePublicData_exp.csv")
        data_pop = pd.read_csv("GooglePublicData_pop.csv")
        data_country = pd.read_csv("GooglePublicData_country.csv")
```

Data Merge in Pandas

Understand the relationship between fertility rate and GDP per capita in 2015

- Plot the following metrics for Year 2015
 - Y: Health / Fertility Rate
 - X: Economic Policy and Debt / GDP per capita (constant 2000 USD)
 - Size: Environment / Population density (people per sq. km of land area)
 - Color: Economic Policy and Debt / Household final consumption expenditure per capita (constant 2000 USD)

1. Extract data for Year 2015

```
In [ ]: data_fer2015 = data_fer[data_fer['year'] == 2015]  
data_fer2015
```

```
In [ ]: data_gdp2015 = data_gdp[data_gdp['year'] == 2015]  
data_gdp2015
```

```
In [ ]: data_exp2015 = data_exp[data_exp['year'] == 2015]  
data_exp2015
```

```
In [ ]: data_pop2015 = data_pop[data_pop['year'] == 2015]  
data_pop2015
```


Data Merge in Pandas

2. Merge different data sets

The four datasets have different numbers of columns and each is for the data in 2015. To combine the four datasets, we can concatenate them with "outer" join.

```
In [ ]: # By default, pd.concat will do row-wise stacking with outer join
data2015 = pd.concat([data_fer2015, data_gdp2015, data_exp2015, data_pop2015])
data2015
```

```
In [ ]: data2015_inner = pd.concat([data_fer2015, data_gdp2015, data_exp2015, data_pop2015], join = "inner")
data2015_inner
```

```
In [ ]: # delete the redundant year column
data2015 = data2015.drop(columns = "year")
data2015
```

```
In [ ]: # By default, pd.concat will do row-wise stacking with outer join
data2015 = pd.concat([data_fer2015, data_gdp2015, data_exp2015, data_pop2015])
data2015
```

```
In [ ]: data2015_inner = pd.concat([data_fer2015, data_gdp2015, data_exp2015, data_pop2015], join = "inner")
data2015_inner
```

```
In [ ]: # delete the redundant year column
data2015 = data2015.drop(columns = "year")
data2015
```

Data Merge in Pandas

The plot function expects each data point on a row, which represents a country's results. Here we just need to do a transpose.

```
In [56]: data2015 = data2015.T  
data2015.head(3)
```

Re-name columns to have meaningful headers

```
In [57]: data2015.columns = ['FertilityRate', 'GDPperCapita', 'HouseholdExpense', 'PopulationDensity']  
data2015.head(3)
```

Data Merge in Pandas

Merge all into one complete dataset

Merge the four imported tables into one such that each subject is one country-year pair, and the four measures are on four separate columns.

| <u>country_id</u> | <u>year</u> | fer | gdp | exp | pop |
|-------------------|-------------|-----|-----|-----|-----|
|-------------------|-------------|-----|-----|-----|-----|

```
In [ ]: data_fer_long = pd.read_csv("GooglePublicData_fer_long.csv")
data_gdp_long = pd.read_csv("GooglePublicData_gdp_long.csv")
data_exp_long = pd.read_csv("GooglePublicData_exp_long.csv")
data_pop_long = pd.read_csv("GooglePublicData_pop_long.csv")
```

```
In [ ]: data_fer_long.head()
```

```
In [ ]: data_gdp_long.head()
```

```
In [ ]: data_all = pd.merge(left=data_fer_long, right=data_gdp_long, on=['country_id', 'year'])
data_all.head()
```

```
In [ ]: data_all = pd.merge(left=data_all, right=data_exp_long, on=['country_id', 'year'])
data_all = pd.merge(left=data_all, right=data_pop_long, on=['country_id', 'year'])
data_all.head()
```

Data Merge in Pandas

Merge all into one complete dataset

Merge the four imported tables into one such that each subject is one country-year pair, and the four measures are on four separate columns.

| <u>country_id</u> | <u>year</u> | fer | gdp | exp | pop |
|-------------------|-------------|-----|-----|-----|-----|
|-------------------|-------------|-----|-----|-----|-----|

```
In [ ]: data_fer_long = pd.read_csv("GooglePublicData_fer_long.csv")
data_gdp_long = pd.read_csv("GooglePublicData_gdp_long.csv")
data_exp_long = pd.read_csv("GooglePublicData_exp_long.csv")
data_pop_long = pd.read_csv("GooglePublicData_pop_long.csv")
```

```
In [ ]: data_fer_long.head()
```

```
In [ ]: data_gdp_long.head()
```

```
In [ ]: data_all = pd.merge(left=data_fer_long, right=data_gdp_long, on=['country_id', 'year'])
data_all.head()
```

```
In [ ]: data_all = pd.merge(left=data_all, right=data_exp_long, on=['country_id', 'year'])
data_all = pd.merge(left=data_all, right=data_pop_long, on=['country_id', 'year'])
data_all.head()
```

Data Merge in Pandas

Merge country info

```
In [66]: data_country.head()
```

Out[66]:

| | country_id | country_name | income_level | latitude | longitude | region |
|---|------------|----------------|--------------|------------|-------------|--------|
| 0 | AFG | Afghanistan | LIC | 33.939110 | 67.709953 | SAS |
| 1 | ALB | Albania | UMC | 41.153332 | 20.168331 | ECS |
| 2 | DZA | Algeria | UMC | 28.033886 | 1.659626 | MEA |
| 3 | ASM | American Samoa | UMC | -14.270972 | -170.132217 | EAS |
| 4 | AND | Andorra | NaN | 42.546245 | 1.601554 | ECS |

```
In [ ]: data_all = pd.merge(left=data_all, right=data_country[['country_id', 'income_level', 'region']], on='country_id')
data_all.head()
```

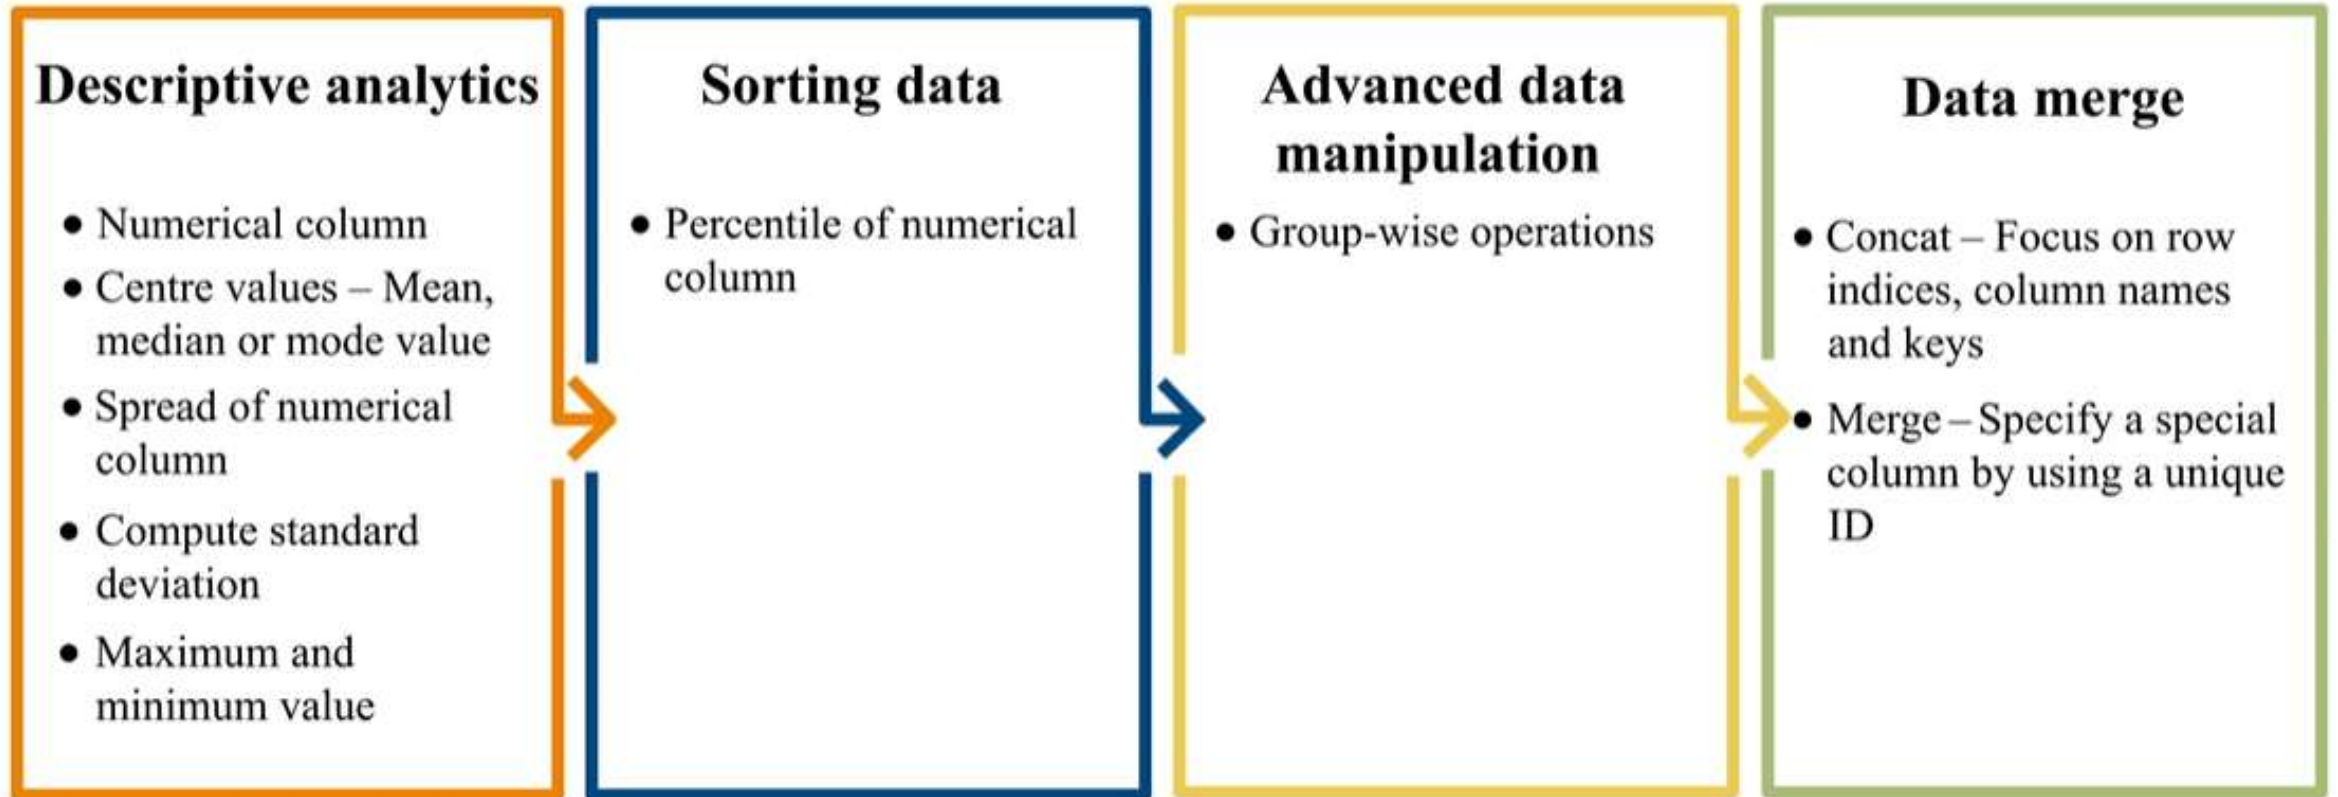
Save the complete dataset for future use

```
In [ ]: data_all.to_csv("GooglePublicData_all.csv", index=False)
```

The background features a large blue triangle on the left side, pointing towards the top-left corner. In the bottom-left corner, there is a small orange triangle pointing towards the bottom-left corner.

Descriptive Analytics with Numerical Summary: Summary

Module Summary





NUS
National University
of Singapore

NUS
BUSINESS
SCHOOL